



Program Optimization – modern CPUs


simon.marechal@synacktiv.com

2023/2024






Amount of cores of Intel consumer CPUs





Base frequency of Intel consumer CPUs



Programming the modern PDP-11



- number of transistors increase exponentially, but processor frequency does not;
- how to increase performance without changing the computing paradigm?
- CPU builders cheat!
 - multiple parallel execution units ;
 - deep pipelining ;
 - and ... eager execution
 - many exploitable bugs!



SYNACKTIV DIGITAL SECURITY 5



CPU architecture

Micro ops



The lowest level programming language available to the programmer is assembly. However, this is a lie!

- assembly instructions are converted to a lower level, undocumented, language: micro ops
- CPUs have several parallel execution ports, virtual registers, etc.
- micro ops can in turn be optimized

(Such knowledge is important for low level optimization, but can be ignored by normal humans)

Execution ports



CMOVcc	r,r	1	1	p06	1	0.5	implicit lock
CMOVcc	r,m	1	2	p06 p23		0.5	
XCHG	r,r	3	3	3p0156	1-2	1	
XCHG	r,m	8	8		21		
XLAT		3	3	p23 2p0156	7	1	not 64 bit
PUSH	r	1	2	p49 p78	3	1	
PUSH	i	1	2	p49 p78		1	
PUSH	m	3	3	p23 p49 p78		1	
PUSH	stack pointer	3	3	p0156 p49 p78		1	
PUSHF(D/Q)		4	4	p06 p1 p49 p78		1	
PUSHA(D)		11	19			8	

Figure 2: Excerpt from

https://www.agner.org/optimize/instruction_tables.pdf

Pipelines



In order to increase throughput, instructions are fed to a pipeline where operations happen in turn:

- ancient, classical, pipeline: instruction fetch, instruction decode, execute, memory access, write back
- Pentium 4 CPUs had a 31-stage pipeline, at the same time AMD64 CPUs had a 12-stage pipeline
- must be kept as full as possible to maintain performance

Out of order execution, register rewriting



Modern CPUs can turn this ...

```
mov eax, [mem1]
imul eax, 6
mov [mem2], eax
mov eax, [mem3]
add eax, 2
mov [mem4], eax
```

... into this!

```
mov eax, [mem1]
mov ebx, [mem3]
add ebx, 2
imul eax, 6
mov [mem2], eax
mov [mem4], ebx
```

Data dependency



```
double list[100];  
double sum = 0.;  
for (int i=0; i<100; i++)  
    sum += list[i];
```

```
double list[100];  
double sumA = 0., sumB = 0.;  
double sumC = 0., sumD = 0.;  
for (int i=0; i<100; i+=4) {  
    sumA += list[i+0];  
    sumB += list[i+1];  
    sumC += list[i+2];  
    sumD += list[i+3];  
}  
double sum = sumA + sumB +  
             sumC + sumD;
```

Branch predictors and eager evaluation



- to maximize performance, pipelines must be kept full
- what happens on a conditional branch?
- CPUs *guess*, and start on a branch
- rollback if necessary
 - can be observable!

Intel guidelines



Static prediction:

- predict unconditional branches to be taken
- predict indirect branches to be NOT taken
- predict backward conditional branches to be taken; rule is suitable for loops
- predict forward conditional branches to be NOT taken

Dynamic prediction:

- undocumented, you are on your own :)



Memory hierarchy

Levels of memory



- registers
- caches
- RAM
- storage
- network?

There is a **huge** performance gap between each category!

Most important consideration!



- a jump is not as bad as an uncached memory read;
- maximizing throughput means avoiding latency: cache aware algorithms, pre-fetching, etc.
- knowledge of micro-ops is pretty important.

Note on *C* is a portable assembler



It is an assembler for a **very** abstract PDP-11.

CPU concepts

- registers
- virtual memory
- execution ports
- cache
- hyper-threading
- SIMD
- ...

C concepts

- variable types
- pointers
- undefined behavior
- sequence points
- aliasing
- struct and union types
- ...

Task



Optimize `simple-c/main.c`

- remove the function calls
- remove useless operations



QUESTIONS?

?



Thank you for your attention

