



Program Optimization – SIMD

simon.marechal@synacktiv.com

2023/2024





What is SIMD?

What is SIMD?



SIMD

single instruction, multiple data

On Intel CPUs



- MMX (1997) with 8, 64 bits, registers (shared with the FP unit)
- SSE (1999) with 8, 128 bits, registers
- SSE2 (2000), SSE3 (2004), SSE4 (2006), incremental updates to SSE
- AVX (2011), with 16, 256 bits, registers
- AVX-512 (2013), with 32, 512 bits, registers

Support



- processors with the *Core* commercial name support these ;
- *Pentium* / *Celeron* do not
- that's why you will need fallback algorithms for older / less capable processors :'(

Compiler support



- most major compiler support the instruction set
- must specify the target architecture though
- usually only works for very simple algorithms



What does it mean

Multiple data in a single register



In a 128 bits register, you can have:

- 2 x 64 bits values (`uint64_t` or `double`) ;
- 4 x 32 bits values (`uint32_t` or `float`) ;
- 8 x 16 bits values (`uint16_t`) ;
- 16 x 8 bits values (`uint8_t`).

And then perform the same operation on all of them at once.

Motivating example



Physics simulation, where, for every *tick*, we do need to update velocity and position:

$$V_t \leftarrow V_{t-1} + A_{t-1} \quad (1)$$

$$P_t \leftarrow P_{t-1} + V_{t-1} \quad (2)$$

In a 3D world, that means 6 additions for every simulation element.

In standard C



```
void add(struct P *l, const struct P *r)
{
    l->x += r->x;
    l->y += r->y;
    l->z += r->z;
    l->w += r->w;
}

...
for (int i = 0; i < OBJECTS; i++)
{
    add(&objects[i].position, &objects[i].velocity);
    add(&objects[i].velocity, &objects[i].acceleration);
}
```

Generated code, -O1, 9.2s



```
0000000000001155 <add>:
    movsd    (%rdi),%xmm0
    addsd    (%rsi),%xmm0
    movsd    %xmm0,(%rdi)
    movsd    0x8(%rdi),%xmm0
    addsd    0x8(%rsi),%xmm0
    movsd    %xmm0,0x8(%rdi)
    movsd    0x10(%rdi),%xmm0
    addsd    0x10(%rsi),%xmm0
    movsd    %xmm0,0x10(%rdi)
    movsd    0x18(%rdi),%xmm0
    addsd    0x18(%rsi),%xmm0
    movsd    %xmm0,0x18(%rdi)
```

Generated code, -O2, 8.2s



Inlined code, operation reordering:

movsd 0x20(%rax),%xmm3	movsd %xmm0,-0x58(%rax)
movsd (%rax),%xmm0	movsd -0x50(%rax),%xmm0
add \$0x60,%rax	movsd %xmm2,-0x38(%rax)
movsd -0x38(%rax),%xmm2	addsd %xmm1,%xmm0
movsd -0x30(%rax),%xmm1	addsd -0x10(%rax),%xmm1
addsd %xmm3,%xmm0	movsd %xmm0,-0x50(%rax)
movsd -0x48(%rax),%xmm4	movsd -0x28(%rax),%xmm0
addsd -0x20(%rax),%xmm3	movsd %xmm1,-0x30(%rax)
movsd %xmm0,-0x60(%rax)	addsd %xmm0,%xmm4
movsd -0x58(%rax),%xmm0	addsd -0x8(%rax),%xmm0
movsd %xmm3,-0x40(%rax)	movsd %xmm4,-0x48(%rax)
addsd %xmm2,%xmm0	movsd %xmm0,-0x28(%rax)
addsd -0x18(%rax),%xmm2	...

Generated code, -O3, 3.9s



Actually use vector instructions:

<code>movupd 0x30(%rdx),%xmm7</code>	<code>addpd %xmm7,%xmm1</code>
<code>movupd 0x10(%rdx),%xmm0</code>	<code>addpd %xmm2,%xmm0</code>
<code>add \$0x60,%rdx</code>	<code>addpd %xmm3,%xmm1</code>
<code>movupd -0x10(%rdx),%xmm2</code>	<code>movups %xmm0,-0x50(%rdx)</code>
<code>movupd -0x40(%rdx),%xmm6</code>	<code>movupd -0x10(%rdx),%xmm0</code>
<code>movupd -0x60(%rdx),%xmm1</code>	<code>movups %xmm1,-0x60(%rdx)</code>
<code>addpd %xmm7,%xmm0</code>	<code>movupd -0x20(%rdx),%xmm1</code>
<code>movupd -0x20(%rdx),%xmm3</code>	<code>addpd %xmm2,%xmm0</code>
<code>addpd %xmm7,%xmm2</code>	<code>addpd %xmm3,%xmm1</code>
<code>movupd -0x40(%rdx),%xmm7</code>	<code>movups %xmm0,-0x30(%rdx)</code>
<code>addpd %xmm6,%xmm3</code>	<code>movups %xmm1,-0x40(%rdx)</code>

Generated code, -O3 -march=native, 3.8s



Wider registers:

```
vmovupd 0x20(%rdx),%ymm0
vmovupd 0x40(%rdx),%ymm1
add      $0x60,%rdx
vaddpd   %ymm0,%ymm1,%ymm2
vaddpd   -0x60(%rdx),%ymm0,%ymm0
vaddpd   %ymm2,%ymm1,%ymm1
vaddpd   %ymm2,%ymm0,%ymm0
vmovupd  %ymm1,-0x40(%rdx)
vmovupd  %ymm0,-0x60(%rdx)
```

Caveat



- Optimizations are VERY fragile!
- most of the time, do not even fire



Task: password cracker

Idea



Adapt your program so that it tries at least 4 plain texts in parallel.

From:


00 40 40 40 40 40 40 40 40 80 00 00 00 00 00 00 00 00
...

To:

00 40 40 40 40 40 40 40 40 40 40 40 40 40 40 40 40 40
00 40 40 40 80 40 40 41 80 40 40 42 80 40 40 43 80
...



QUESTIONS?



Thank you for your attention

