



Program Optimization – GPUs

simon.marechal@synacktiv.com

2023/2024





Introduction



History



graphics cards are really good at parallel processing,

- early efforts for general purpose computation with shaders
- 2006, CUDA
- 2009, OpenCL
 - also for CPUs, FPGAs, etc.
- 2010, (old) oclHashcat, very limited
- 2011, oclHashcat-plus
- 2013, oclHashcat



When to use GPUs?



- very parallel workloads (same operations on a large amount of data),
- ideally, cut in a lot of parts,
- high arithmetic density (ratio of calculation over memory operations)



Difference with CPUs



- much, *much* simpler execution units
 - in order execution (except RTX 4xxx series?)
 - no branch prediction
- many threads in parallel (1k / 2k)



OpenCL programming concepts



Platform model



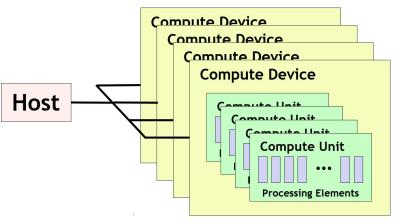


Figure 1: from Khronos Group OpenCL Guide



Execution model



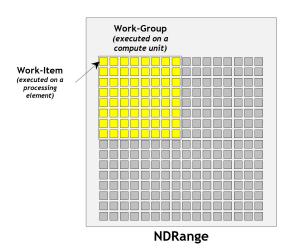


Figure 2: from Khronos Group OpenCL Guide



Memory model



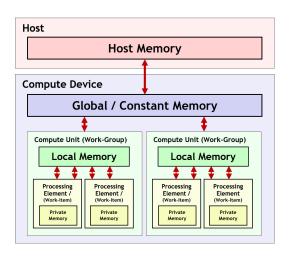


Figure 3: from Khronos Group OpenCL Guide



OpenCL C



- looks a lot like C, but isn't (we will not study how it differs)
- specific concepts:
 - address spaces, local, global, private, constant
 - work item / work group
 - command queues
- get_global_id()

Task - GPU cracker



- you are given an initial C adapter program
- write the .cl file!
 - simple cracker
 - optimize it
 - implement better queuing

Why are GPUs fast?



- many, many threads can run on a single processing unit
- they are scheduled so as to mask memory latencies
- this means it is easier to approach peak execution rate for those processors









