



Performance Optimization – general guidelines

simon.marechal@synacktiv.com

2021/2022





Introduction

Start with the basics



Make it work, make it right, make it fast.

– Kent Beck

Better than the famous Dijkstra quote!

- get a baseline that you are confident will work
- find a way to check there are no regressions during optimization
 - when working on someone else's code, it is easy to miss invariants
 - your future self is someone else :)

Architecture and system design



It is important not to optimize too much during design

- makes refactoring, changes, more costly
- wasted work

However, it is important to keep performance in mind!

- architectural choices
- it is good to have an feeling on how fast it will be in the end
- it is OK to work in a high level language, as long as you can get low level

Measure first



Always measure before optimizing!

- the compiler might have done a better job already;
- the hot path might not be where you expect it;
- the algorithmically sensible decision might turn out to be a disaster in practice (small n , high constant factors)

Make the slow parts faster :)



- go for the parts that eats up most of the computation time first
 - however that might be normal
 - that is why you should have an idea on how fast things should be first



Minimal tool set for program optimization

Micro-benchmarks



A test suite that measure the speed of a small part of your program

- kind of artificial, high variance between runs is expected
 - 5-15% on a typical computer
- the test framework must:
 - warm-up the CPU
 - make sure the algorithm is not optimized away
 - display distributions
- write benches with several input sizes
- [Click here for a good example](#)

Execution traces



- realistic measurements
- however, many factors can muddle the conclusions
 - dependencies, such as storage, databases, etc.
 - can be hard to interpret





Other concerns

Talking with the kernel



System calls are very expensive!

- reduce them (ie. techniques for caching the current time)
- zero-copy tools
- batch them (io-uring techniques)

VMs and containers impose a tax on everything, although it tends to be less and less expensive


Note that most large orgs have a dedicated kernel team to just *make it work*

- super common systems like TCP are not tuned anymore

Hypervisor performance by generation



VM Improvements



VM: Virtual Machine. HW: Hardware.

VS: Virt. in software. VH: Virt. in hardware. P: Paravirt. Not all combinations shown.

SR-IOV(net): igbena driver. SR-IOV(storage): nvme driver.

<http://www.brendangregg.com/blog/2017-11-29/aws-ec2-virtualization-2017.html>

LISA'21 Computing Performance: On the Horizon (Brendan Gregg)

Source:
[Gregg 17]

Local system performance



- for storage dependent elements, use the fastest storage you can afford and **never** use RAID5
- scheduling makes all the difference ...
- the performance of many systems depend on a centralized persistent store, like a database
 - learn how to optimize queries, table representation, indexes
 - investigate all the tunable settings (database and kernel)
 - if at all possible, avoid virtualizing it!


Whole system performance



- everything must be tuned
- observability is *really* hard
- look at everything Brendan Gregg does



QUESTIONS?



Thank you for your attention

